

Frill — Functional Language for Z80

Frill is an ML-style functional language that compiles to Z80 assembly through the MinZ compiler pipeline. The name is ironic — “frills” (ruffles, decorations) for the most constrained hardware imaginable.

Philosophy

“What if OCaml/F# ran on a ZX Spectrum?”

Frill takes the best ideas from ML-family languages and strips them down to what makes sense on 8-bit hardware with 64KB RAM:

From	Idea	Frill adaptation
OCaml/F#	ADT, pattern matching, HM inference	ADT as tagged u8, match as if-chain
F#	Pipe operator <code>\ </code> >	Direct — zero overhead
Haskell	Where-clauses, type classes	let-in (where = future)
Lean 4	Strict, RC + destructive reuse	Strict by default, stack-based lifetime
Idris 2	Linear types (QTT)	Future: 1 x = use exactly once
Rust	Ownership without borrow	Implicit through linearity (planned)

From	Idea	Frill adaptation
	checker	
	noise	

Three design principles:

1. **Strict** — no thunks, no lazy evaluation, no GC needed
2. **Size-known** — every type has compile-time-known size (u8, u16, tagged unions)
3. **Zero-cost** — functional abstractions compile to the same Z80 code a hand-written assembly programmer would write

Quick Start

(* hello.frl — your first Frill program *)

```
let double (x : u8) : u8 = x + x
let inc (x : u8) : u8 = x + 1
```

```
(* pipe: 3 -> double -> inc = 7 *)
let transform (x : u8) : u8 = x |> double |> inc
```

```
assert transform 3 == 7
assert transform 10 == 21
```

Compile and verify:

```
minzc hello.frl # -> hello.a80 (Z80 assembly)
minzc hello.frl -o hello.bin # -> binary
```

Asserts are verified at compile time via the MIR2 VM — no runtime overhead.

Syntax Reference

Functions

ML-style: let name (param : type) ... : return_type = body

```
let add (a : u8) (b : u8) : u8 = a + b
let square (x : u8) : u8 = x * x
let negate (x : i8) : i8 = 0 - x
```

Compiles to:

```
add:      ADD A, C / RET      ; 2 instructions, 8 T-states
square:   JP __mul8          ; tail call
```

Types

Fixed-size, known at compile time:

Type	Size	Range
u8	1 byte	0..255
u16	2 bytes	0..65535
i8	1 byte	-128..127
i16	2 bytes	-32768..32767
bool	1 byte	true/false

No heap. No GC. No boxing. Ever.

If-Then-Else (Expression)

```
let max (a : u8) (b : u8) : u8 =  
  if a > b then a else b
```

```
let clamp (x : u8) (lo : u8) (hi : u8) : u8 =  
  if x < lo then lo  
  else if x > hi then hi  
  else x
```

If is an **expression** — always returns a value, always has both branches.

Let-In (Scoped Bindings)

```
let distance (x1 : u8) (y1 : u8) (x2 : u8) (y2 : u8) : u8 =  
  let dx = abs_diff x1 x2 in  
  let dy = abs_diff y1 y2 in  
  dx + dy
```

Let-in chains are desugared to HIR VarDeclStmt sequences — zero overhead, no closures, no heap allocation. Variables live on the Z80 stack or in registers.

Pipe Operator |>

```
let process (x : u8) : u8 = x |> double |> inc |> square
```

```
(* equivalent to: square(inc(double(x))) *)
```

Pipes with extra arguments:

```
let offset (x : u8) : u8 = x |> add 10
```

```
(* equivalent to: add(x, 10) *)
```

The pipe operator is purely syntactic — it rearranges the call, generating identical code to the nested version.

Algebraic Data Types (ADT)

```
type Direction = North | South | East | West
```

```
type Color = Red | Green | Blue | Yellow | Cyan | Magenta | White  
| Black
```

Constructors compile to sequential integer tags (North=0, South=1, ...). An ADT value is a single u8 — the tag. No heap, no pointers, no vtables.

Constructors work as expressions:

```
let default_dir (x : u8) : u8 = North      (* returns 0 *)  
let go_east (x : u8) : u8 = East          (* returns 2 *)
```

Pattern Matching

```
let describe (d : u8) : u8 =  
  match d with  
  | North -> 0  
  | South -> 1  
  | East  -> 2  
  | West  -> 3  
  | _     -> 99  
end
```

```
let is_weekend (d : u8) : u8 =  
  match d with  
  | Sat -> 1  
  | Sun -> 1  
  | _   -> 0  
end
```

match compiles to a chain of comparisons — equivalent to a switch/case on the tag value. The _ wildcard is the default/else branch.

On Z80 this becomes:

```
CP 5          ; == Sat?  
JR Z, .yes  
CP 6          ; == Sun?  
JR Z, .yes  
XOR A         ; default: 0  
RET  
.yes:  
LD A, 1  
RET
```

Assertions (Compile-Time Tests)

```
assert add 3 5 == 8
assert max 10 20 == 20
assert is_weekend 5 == 1
```

Asserts run on the MIR2 virtual machine at compile time. They are NOT included in the binary — zero runtime cost. If an assert fails, compilation stops with an error showing expected vs actual.

Comments

```
(* This is a block comment – OCaml style *)
-- This is a line comment – Haskell/SQL style
```

How It Works: The Pipeline

```
Frill source (.frl)
  |
  v
[Frill Parser] ---- pkg/frill/frill.go (809 LOC)
  |                    Lexer + recursive descent parser
  |                    Desugars: pipe, let-in, match, ADT
  v
[HIR Module] ---- pkg/hir/hir.go
  |                    Typed AST: Func, VarDecl, If, Return,
Call, BinExpr...
  |                    Target-neutral, structured control flow
  v
[MIR2 Module] ---- pkg/hir/lower.go → pkg/mir2/
  |                    SSA form, basic blocks, virtual registers
  |                    Optimization: constant fold, DSE, inline,
PBQP alloc
  v
[LIR Backend] ---- pkg/lir/
  |                    ISLE instruction combining + WFC register
allocation
  |                    Pattern-driven: data tables, not code
  v
[Z80 Assembly] ---- .a80 file
  |
  v
[MZA Assembler] ---- pkg/z80asm/
  |                    Table-driven encoder, 1335/1335 FUSE tests
  v
[Binary] ---- .bin / .sna / .tap / .com
```

Each Frill construct maps to HIR nodes:

Frill	HIR
let f (x : u8) : u8 = body	Func{Name, Params, RetTy, Body}
x + y	BinExpr{"+", VarRef{x}, VarRef{y}}
if c then a else b	CondExpr{c, a, b}
let x = e in body	VarDeclStmt{x, e} + body
= body where x = e	VarDeclStmt{x, e} before body
x \> f	CallExpr{f, [x]}
x \> \n\ n + 1	Lambda → auto-generated __lambda_N function
match x with \ P -> e end	Nested CondExpr chain
type T = A \ B	Constructor registry (tag integers)
type T = A \ B of u8	Tag + payload encoded as u16
"hello\n"	CString interned in module, AddrOfExpr
assert f x == n	Assert{f, [x], n, via:"mir2"}

What's Implemented (30 features)

Core Language



Function definitions with typed parameters



Return type inference (let double (x : u8) = x + x)



Arithmetic: + - * / %



Comparisons: == != < <= > >=



If-then-else expressions (nested: if/else if/else)



Unary minus (-x), hex literals (0xFF), booleans (true/false)



String literals with escapes ("hello\n")

Functional Features



Let-in bindings with chaining (let x = e1 in let y = e2 in
body)



Where-clauses (= body where x = e)



Pipe operator `|>` with extra args (`x |> add 10`)



Function composition `>>` (`let transform = double >> inc >> square`)



Lambda in pipes (`x |> |n| n + 1`)



Currying / partial application (`let inc = add 1`)



Operator sections (`x |> (+10)`)



Recursion and mutual recursion (factorial, fibonacci, GCD)

Type System



Algebraic data types (`type Color = Red | Green | Blue`)



ADT with payload (`type Option = None | Some of u8`)



Pattern matching with `match/with/end`



Exhaustive match checking (compile error: missing `Blue`)



Guards in match (`| _ when n > 10 -> ...`)



Record types (`type Point = { x : u8, y : u8 }`)



Tuples — multi-return (`u8, u8`) + destructuring `let (q, r) = divmod x y in`

Verification & Testing



Compile-time assertions (`assert gcd 12 8 == 4`)



Property-based testing (`prop |x| double x == x + x` — tests all 256 `u8` values)



MIR2 VM verification — zero runtime cost

Modularity & I/O



Import `.frl` files (`import "math.frl"`)

- Cross-language import .nanz files (import "io_cpm.nanz")
- Extern declarations (extern putchar (ch : u8) : void)
- Inline assembly (asm "LD E, A / CALL 5")
- Do notation for side effects (do putchar 72)
- CLI: minzc program.frl / MZV: mzv program.frl

Stats

- ~1500 LOC parser (pkg/frill/frill.go)
- 13 unit tests
- 6 demo programs (basics, math, game, showcase, stdlib_demo, hello_cpm)
- 26-function math stdlib
- 100+ compile-time assertions across demos
- 768+ property checks (3 props x 256 values)
- CP/M hello world: 720 bytes, prints text on real Z80 emulator

Benchmark: Frill vs SDCC (C)

GCD (Euclid's Algorithm)

Frill source (1 line):

```
let gcd (a : u8) (b : u8) = if b == 0 then a else gcd b (a % b)
```

C source (SDCC):

```
unsigned char gcd(unsigned char a, unsigned char b) {
    if (b == 0) return a;
    return gcd(b, a % b);
}
```

Metric	Frill/MinZ	SDCC/C
Self-contained	Yes (inline div8)	No (__moduchar library)
Total instructions	23	12 + ~30 (library)
Tail call optimized	Yes (JP)	Yes (JR)
T-states/iteration	~342T	~294T
External dependencies	0	1 runtime function

Both generate tail-recursive code. MinZ inlines the division loop (no runtime library), SDCC calls `__moduchar`. Total code footprint favors MinZ.

Simple Functions

```
let add (a : u8) (b : u8) = a + b      →  ADD A, C / RET      (2
insts)
let double (x : u8) = x + x            →  ADD A, A / RET      (2
insts)
let max (a : u8) (b : u8) = ...        →  CP B / RET          (2
insts)
```

Optimal — matches hand-written assembly.

What's Missing (Roadmap)

Frill-1: Full ML (mostly done)



Tuples — `let divmod (a : u8) (b : u8) : (u8, u8) = (a / b, a % b)`



Nested function calls — `f (g x)` works via parens



Pattern match on payload — `| Some x -> x + 1` (parses, runtime u16 bug)

Frill-2: Advanced Types (in progress)



Type classes — `class Show where show : u8 + instance Show Color where ...`

- Compiled via monomorphization (mangled names: `show_u8`, `show_Color`)
- No vtables, no runtime dispatch — zero cost on Z80



Parametric polymorphism — `let id (x : 'a) : 'a = x`

- Monomorphized at call site (like C++ templates, no runtime cost)

Frill-3: Linear Types & TSMC Spill (in progress)



Linearity analysis — compiler classifies each param as 0/1/ ω uses

- 0 (erased): param unused → dead param elimination

- 1 (linear): used once → register freed after use, no save needed
- ω (shared): used 2+ times → needs preservation



Linear type annotations — `let consume (! buf : u8) = ...`
(must use once)



Erased types — `let size (~ T : Type) : u16` — type-level only, zero runtime



TSMC spill slots — self-modifying code for register preservation (see below)



Dependent types — `Vec (n : u8) (a : Type)` — length-indexed vectors

TSMC Tunnels: Linearity Meets Self-Modifying Code

On Z80, the traditional way to save a register across a function call is PUSH/POP (21 T-states, uses stack). Frill's linearity analysis enables a better approach: **TSMC (True Self-Modifying Code) spill slots**.

A TSMC tunnel saves a value by patching it into an instruction's immediate field, then reloads it later by executing that instruction:

```
; Save A into TSMC slot (13T)
    LD (_tsmc_slot_a1), A

; ... function call or complex code that clobbers A ...

; Reload A from TSMC slot (7T) – the immediate byte was patched
above
.reload:
    LD A, 0 ; ← this 0 gets overwritten to the
saved value
_tsmc_slot_a1 EQU .reload + 1 ; points at the immediate byte
```

Cost comparison:

Method	Save	Reload	Total	Stack?
PUSH/POP	11T	10T	21T	Yes (2 bytes)
TSMC tunnel	13T	7T	20T	No
Shadow (EX AF,AF')	4T	4T	8T	No (but only 1 slot)
IXH/IXL (eZ80)	8T	8T	16T	No (2 slots)

Method	Save	Reload	Total	Stack?
Linear (1 use)	0T	0T	0T	No spill needed!

For non-A registers, the TSMC save goes through A:

```
; Save D into TSMC slot
    EX AF, AF'          ; save current A
    LD A, D              ; move D → A
    LD (_tsmc_slot_d), A ; patch the reload instruction
    EX AF, AF'          ; restore A

; Reload D from TSMC slot
.reload_d:
    LD D, 0              ; ← patched with saved value
_tsmc_slot_d EQU .reload_d + 1
```

Why TSMC tunnels are recursion-friendly: Unlike stack spills, TSMC slots don't grow with recursion depth. Each function has a fixed set of TSMC slots in the data section. For non-recursive functions this is always correct. For recursive functions, TSMC slots work when there's no recursive call between the save (tunnel start) and reload (tunnel end) — the compiler verifies this statically via the linearity analysis.

Connection to QTT: A linear parameter (quantity 1) never needs a TSMC tunnel — it's used once and the register is freed. An erased parameter (quantity 0) never needs a register at all. Only shared parameters (quantity ω) might need TSMC tunnels, and the compiler knows exactly which ones at compile time.

Why These Matter on Z80

Linear types formalize what Z80 programmers do intuitively. When you pass a buffer to a function and never use it again, the compiler knows it can reuse the memory. No GC, no RC overhead — the stack frame IS the allocator.

Dependent types catch buffer overflows at compile time. `Vec 10 u8` guarantees exactly 10 elements — the compiler rejects `index 11`.

Type classes replace vtables with monomorphization. `print x` where `x : u8` compiles to `CALL print_u8` — zero indirection.

Examples

Collatz Conjecture

```
let collatz_step (n : u8) : u8 =  
  if n == 1 then 1  
  else if n % 2 == 0 then n / 2  
  else n * 3 + 1
```

```
assert collatz_step 6 == 3  
assert collatz_step 5 == 16  
assert collatz_step 1 == 1
```

Day of Week

```
type Day = Mon | Tue | Wed | Thu | Fri | Sat | Sun
```

```
let is_weekend (d : u8) : u8 =  
  match d with  
  | Sat -> 1  
  | Sun -> 1  
  | _ -> 0  
end
```

```
assert is_weekend 5 == 1      (* Sat *)  
assert is_weekend 0 == 0      (* Mon *)
```

Pipeline Processing

```
let double (x : u8) : u8 = x + x  
let inc (x : u8) : u8 = x + 1  
let square (x : u8) : u8 = x * x
```

```
let process (x : u8) : u8 = x |> double |> inc |> square
```

```
assert process 3 == 49      (* 3 -> 6 -> 7 -> 49 *)
```

Recursion

```
let factorial (n : u8) : u8 =  
  if n == 0 then 1  
  else n * factorial (n - 1)
```

```
let gcd (a : u8) (b : u8) : u8 =  
  if b == 0 then a  
  else gcd b (a % b)
```

```
assert factorial 5 == 120
assert gcd 12 8 == 4
```

Lambda in Pipes

```
let process (x : u8) : u8 =
  x |> |n| n + 1 |> |n| n * 2 |> |n| n - 3

assert process 5 == 9  (* 5 -> 6 -> 12 -> 9 *)
```

Where-Clause

```
let hyp2 (x : u8) (y : u8) : u8 = xx + yy
  where xx = x * x
  where yy = y * y

assert hyp2 3 4 == 25
```

Let-In + Pipe

```
let compute (x : u8) : u8 =
  let a = x + 1 in
  let b = a * 2 in
  let c = b - x in
  c |> |n| n + n

assert compute 4 == 12  (* a=5, b=10, c=6, double=12 *)
```

ADT with Payload

```
type Option = None | Some of u8

let is_some (x : u8) : u8 = __tag (Some x)
let value (x : u8) : u8 = __payload (Some x)

assert is_some 42 == 1
assert value 42 == 42
```

Exhaustive Match

```
type Color = Red | Green | Blue

(* This compiles: *)
let name (c : u8) : u8 =
  match c with | Red -> 0 | Green -> 1 | Blue -> 2 end

(* This is a compile error: "non-exhaustive match on Color:
missing Blue" *)
```

```
(* let bad (c : u8) : u8 = match c with | Red -> 0 | Green -> 1
end *)
```

Generated Assembly — Three Backends Compared

MinZ has three codegen backends. Same Frill source, three different outputs.

add — let add (a : u8) (b : u8) = a + b

```
; LIR (ISLE+WFC) — optimal
add:
    ADD A, C
    RET
```

```
; PBQP (MIR2) — identical
add:
    ADD A, C
    RET
```

Both backends generate the same optimal 2-instruction code: a=param1 in A, b=param2 in C. ADD A,C. Return result in A. Can't do better.

double — let double (x : u8) = x + x

```
; Both backends:
double:
    ADD A, A          ; x + x = shift left, 4T
    RET
```

Self-add instead of SLA (which would also work). Both backends identical.

abs_diff — let abs_diff (a : u8) (b : u8) = if a > b then a - b else b - a

```
; PBQP — uses conditional return (TermCondRet)
abs_diff:
    SUB C          ; A = a - b
    RET NC         ; if no carry (a >= b), return a-b
    NEG           ; else negate: A = b - a
    RET           ; 4 instructions, branchless-ish
```

```
; LIR — different approach
abs_diff:
    SUB B          ; subtract
```

```

CP B          ; compare
RET          ; (simplified, may lose accuracy)

```

PBQP generates elegant code: subtract, check carry flag, conditionally negate. Single compare + conditional return — classic Z80 idiom.

factorial — **let factorial (n : u8) = if n == 0 then 1 else n * factorial (n - 1)**

```

; PBQP – recursive with tail-call to __mul8
factorial:
    AND A          ; test n == 0
    JRS NZ, .else  ; if not zero, recurse
    LD A, 1        ; base case: return 1
    RET
.else:
    LD C, 1
    SUB 1          ; n - 1
    LD B, A
    CALL factorial ; recursive call
    CALL __mul8    ; n * factorial(n-1)
    RET

```

Recursive factorial: 11 instructions + shared __mul8 runtime. CALL/RET stack manages the recursion naturally on Z80.

gcd —

let gcd (a : u8) (b : u8) = if b == 0 then a else gcd b (a % b)

```

; PBQP – tail-recursive Euclid with inline division
gcd:
    LD A, C
    AND A          ; test b == 0
    JRS NZ, .else
    RET            ; base: return a (already in A)
.else:
    LD B, C        ; setup for div8
    XOR A
    LD D, 8
.div8:
    SLA B          ; shift-and-subtract division
    RLA
    CP C
    JR C, .skip
    SUB C
    INC B
.skip:
    DEC D

```

```

JR NZ, .div8          ; 8 iterations
LD C, A                ; remainder → b
LD A, B                ; quotient (unused, but b is new a)
JP gcd                 ; tail call – no stack growth!

```

Euclid's algorithm: 23 instructions, fully self-contained. Division inlined as 8-cycle shift-and-subtract loop. **Tail call** (JP gcd instead of CALL+RET) — zero stack growth for any input.

inc — let inc (x : u8) = x + 1

```

; PBQP – optimal
inc:
    INC A                ; 4T, 1 byte – can't do better
    RET

; LIR – one extra instruction
inc:
    LD C, 1              ; materialise constant (unnecessary)
    INC A
    RET

```

PBQP wins here — LIR materialises the constant 1 even though INC A doesn't need it. A future ISLE peephole rule could eliminate this.

square — let square (x : u8) = x * x

```

; LIR – tail call to runtime
square:
    JP __mul8             ; A=x already, B=x (alias), tail call

; PBQP – explicit setup
square:
    LD B, A                ; B = x (for __mul8: A * B → A)
    CALL __mul8
    RET

```

LIR uses a tail call (JP instead of CALL+RET), saving 17 T-states. Both use the shared __mul8 runtime (~80T software multiply). On eZ80, this would be a single MLT instruction (6T).

eZ80 ADL — let add (a : u8) (b : u8) = a + b

```

; eZ80 ADL mode – same optimal code, 24-bit addresses
    .ASSUME ADL=1
    ORG $040045
add:

```



```
ADD A, C
RET
```

eZ80 backend wraps Z80 assembly with `.ASSUME ADL=1` header. MZA assembler handles 24-bit address encoding automatically. Same code, 3-byte immediates instead of 2-byte.

Honest Assessment: Code Quality

What's Good

Simple functions are optimal. `add`, `double`, `inc` generate the exact same code a human would write. `ADD A, C / RET` — you literally cannot do better. The PBQP allocator gets register assignment right for leaf functions.

Conditional returns are clever. `abs_diff` generates `SUB C / RET NC / NEG / RET` — four instructions, no branch. The MIR2 `TermCondRet` produces idiomatic Z80 that experienced asm programmers would recognize.

Tail call optimization works. `gcd` compiles to `JP gcd` instead of `CALL gcd / RET`. Zero stack growth for any recursion depth. This is critical on Z80 where the stack is typically <256 bytes.

Division is self-contained. No runtime library needed. The 8-bit `div` loop is standard shift-and-subtract — correct if not fancy.

What's Not Good

LIR has bugs with recursive functions. `gcd` through LIR generates `JP gcd` (infinite loop) — the constant folder incorrectly eliminates the conditional branch. PBQP path works; LIR does not for recursion. This is why `--lir=false` is required for recursive Frill programs.

LIR materialises unnecessary constants. `inc` generates `LD C, 1 / INC A / RET` instead of just `INC A / RET`. The constant 1 is loaded into C but never used. A peephole rule should catch this, but doesn't yet.

8-bit multiply is expensive. `square` calls `__mul8` (~80 T-states, ~30 bytes). On eZ80, `MLT BC` would be 6 T-states / 2 bytes. The Z80 has no hardware multiply, so this is unavoidable, but the `__mul8` routine itself could be smaller with loop unrolling or Karatsuba for known-small operands.

Spill data is always emitted. Every compiled file includes `mem0: DW 0` through `mem3: DW 0` and `tsmc0..7` — 24 bytes of spill slots even when unused. This should be emitted only when needed.

Register pressure in complex expressions. `print_u8(fib 7)` fails because the return value from `fib` gets clobbered before being passed to `print_u8`. The Z80 has only 7 general-purpose 8-bit registers — complex expression trees exceed this and need spill/reload. The PBQP allocator handles this correctly for most cases, but deeply nested calls at the Frill level can produce incorrect code because the HIR→MIR2 lowering doesn't insert explicit temporaries for intermediate values.

No inlining across Frill functions. `x |> double |> inc` generates `CALL double / CALL inc` — two function calls (34 T-states overhead). Hand-written: `ADD A, A / INC A` — zero overhead. The MIR2 `InlineTrivial` pass could inline these, but it requires the functions to be in the same module and small enough. Currently works for some cases but not reliably.

Compared to SDCC (C Compiler)

Aspect	MinZ/Frill	SDCC
Simple functions	Optimal (same as hand-asm)	Optimal
Register allocation	PBQP — good for leaf, struggles with deep nesting	Graph coloring — more mature
Tail calls	Yes (JP)	Yes (JR)
Division	Inline loop (self-contained)	Library call (smaller caller)
Multiply	<code>__mul8</code> shared routine	<code>__moduchar</code> library
Inlining	Limited (InlineTrivial)	Aggressive (-O2)
Code size overhead	~24 bytes spill slots	~0 bytes (only what's used)
Maturity	~6 months, experimental	~20 years, production

Bottom line: Frill generates correct, sometimes optimal Z80 code for simple to medium functions. Complex programs with deep recursion or many intermediate values hit register pressure

limits. The sweet spot is the same as ML's sweet spot: small, pure, composable functions — which happen to be exactly what the Z80's register file can handle.

Design Decisions

Why strict? Lazy evaluation requires thunks (heap-allocated closures). On Z80 with 48KB usable RAM and no MMU, heap allocation means manual memory management or GC — both unacceptable. Strict evaluation uses the stack naturally.

Why no GC? The Z80 runs at 3.5 MHz. A GC pause of even 1000 cycles (0.3ms) is visible in games running at 50fps. Linear types + stack allocation give deterministic performance.

Why tags as u8? An ADT with ≤ 256 constructors fits in a single byte. Pattern matching becomes a CP instruction (4 T-states). Tagged unions with payload will use 1 byte tag + N bytes payload — still fixed-size, still stack-allocated.

Why desugar match to if-chain? For ≤ 8 arms, a chain of CP/JR Z instructions is faster than a jump table on Z80 (no multiply for index, no memory load for target). The compiler can switch to jump tables for larger matches in the future.

Relationship to Other MinZ Languages

MinZ has 8 frontend languages, all sharing the same backend:

```
Frill (.frl)  ----\  
Nanz (.nanz) ----+  
C89 (.c)     ----+  
Pascal (.pas) ----+----> HIR ---> MIR2 ---> LIR ---> Z80/eZ80  
Lanz (.lanz) ----+  
Lizp (.lizp) ----+  
PL/M (.plm)  ----+  
ABAP (.abap) ----/
```

Frill is the **functional** member of the family. Use it when you want: - Pattern matching on structured data - Pipeline-style data transformation - Compile-time verified correctness (assert) - Clean, expression-oriented code

Use Nanz when you need: - Mutable state, while loops, imperative I/O - Inline assembly (`asm z80 { ... }`) - Stdlib imports (TUI, graphics, sound) - Closures and iterators